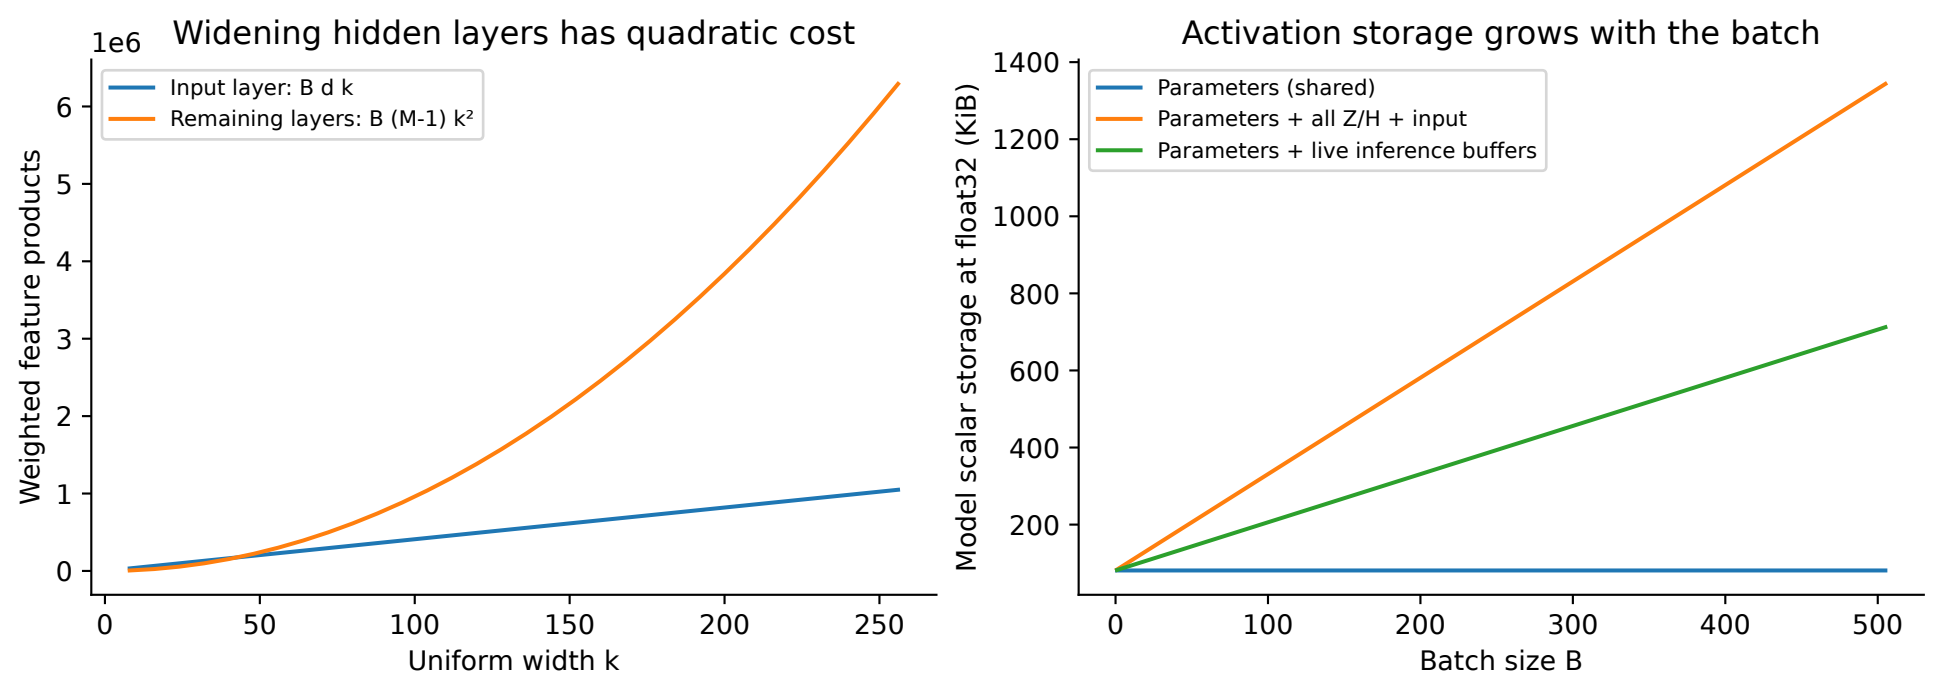
Widening hidden layers has quadratic cost
1e6
Input layer: B d k
Remaining layers: B (M-1) k²
Weighted feature products
Uniform width k
Activation storage grows with the batch
Parameters (shared)
Parameters + all Z/H + input
Parameters + live inference buffers
Model scalar storage at float32 (KiB)
Batch size B